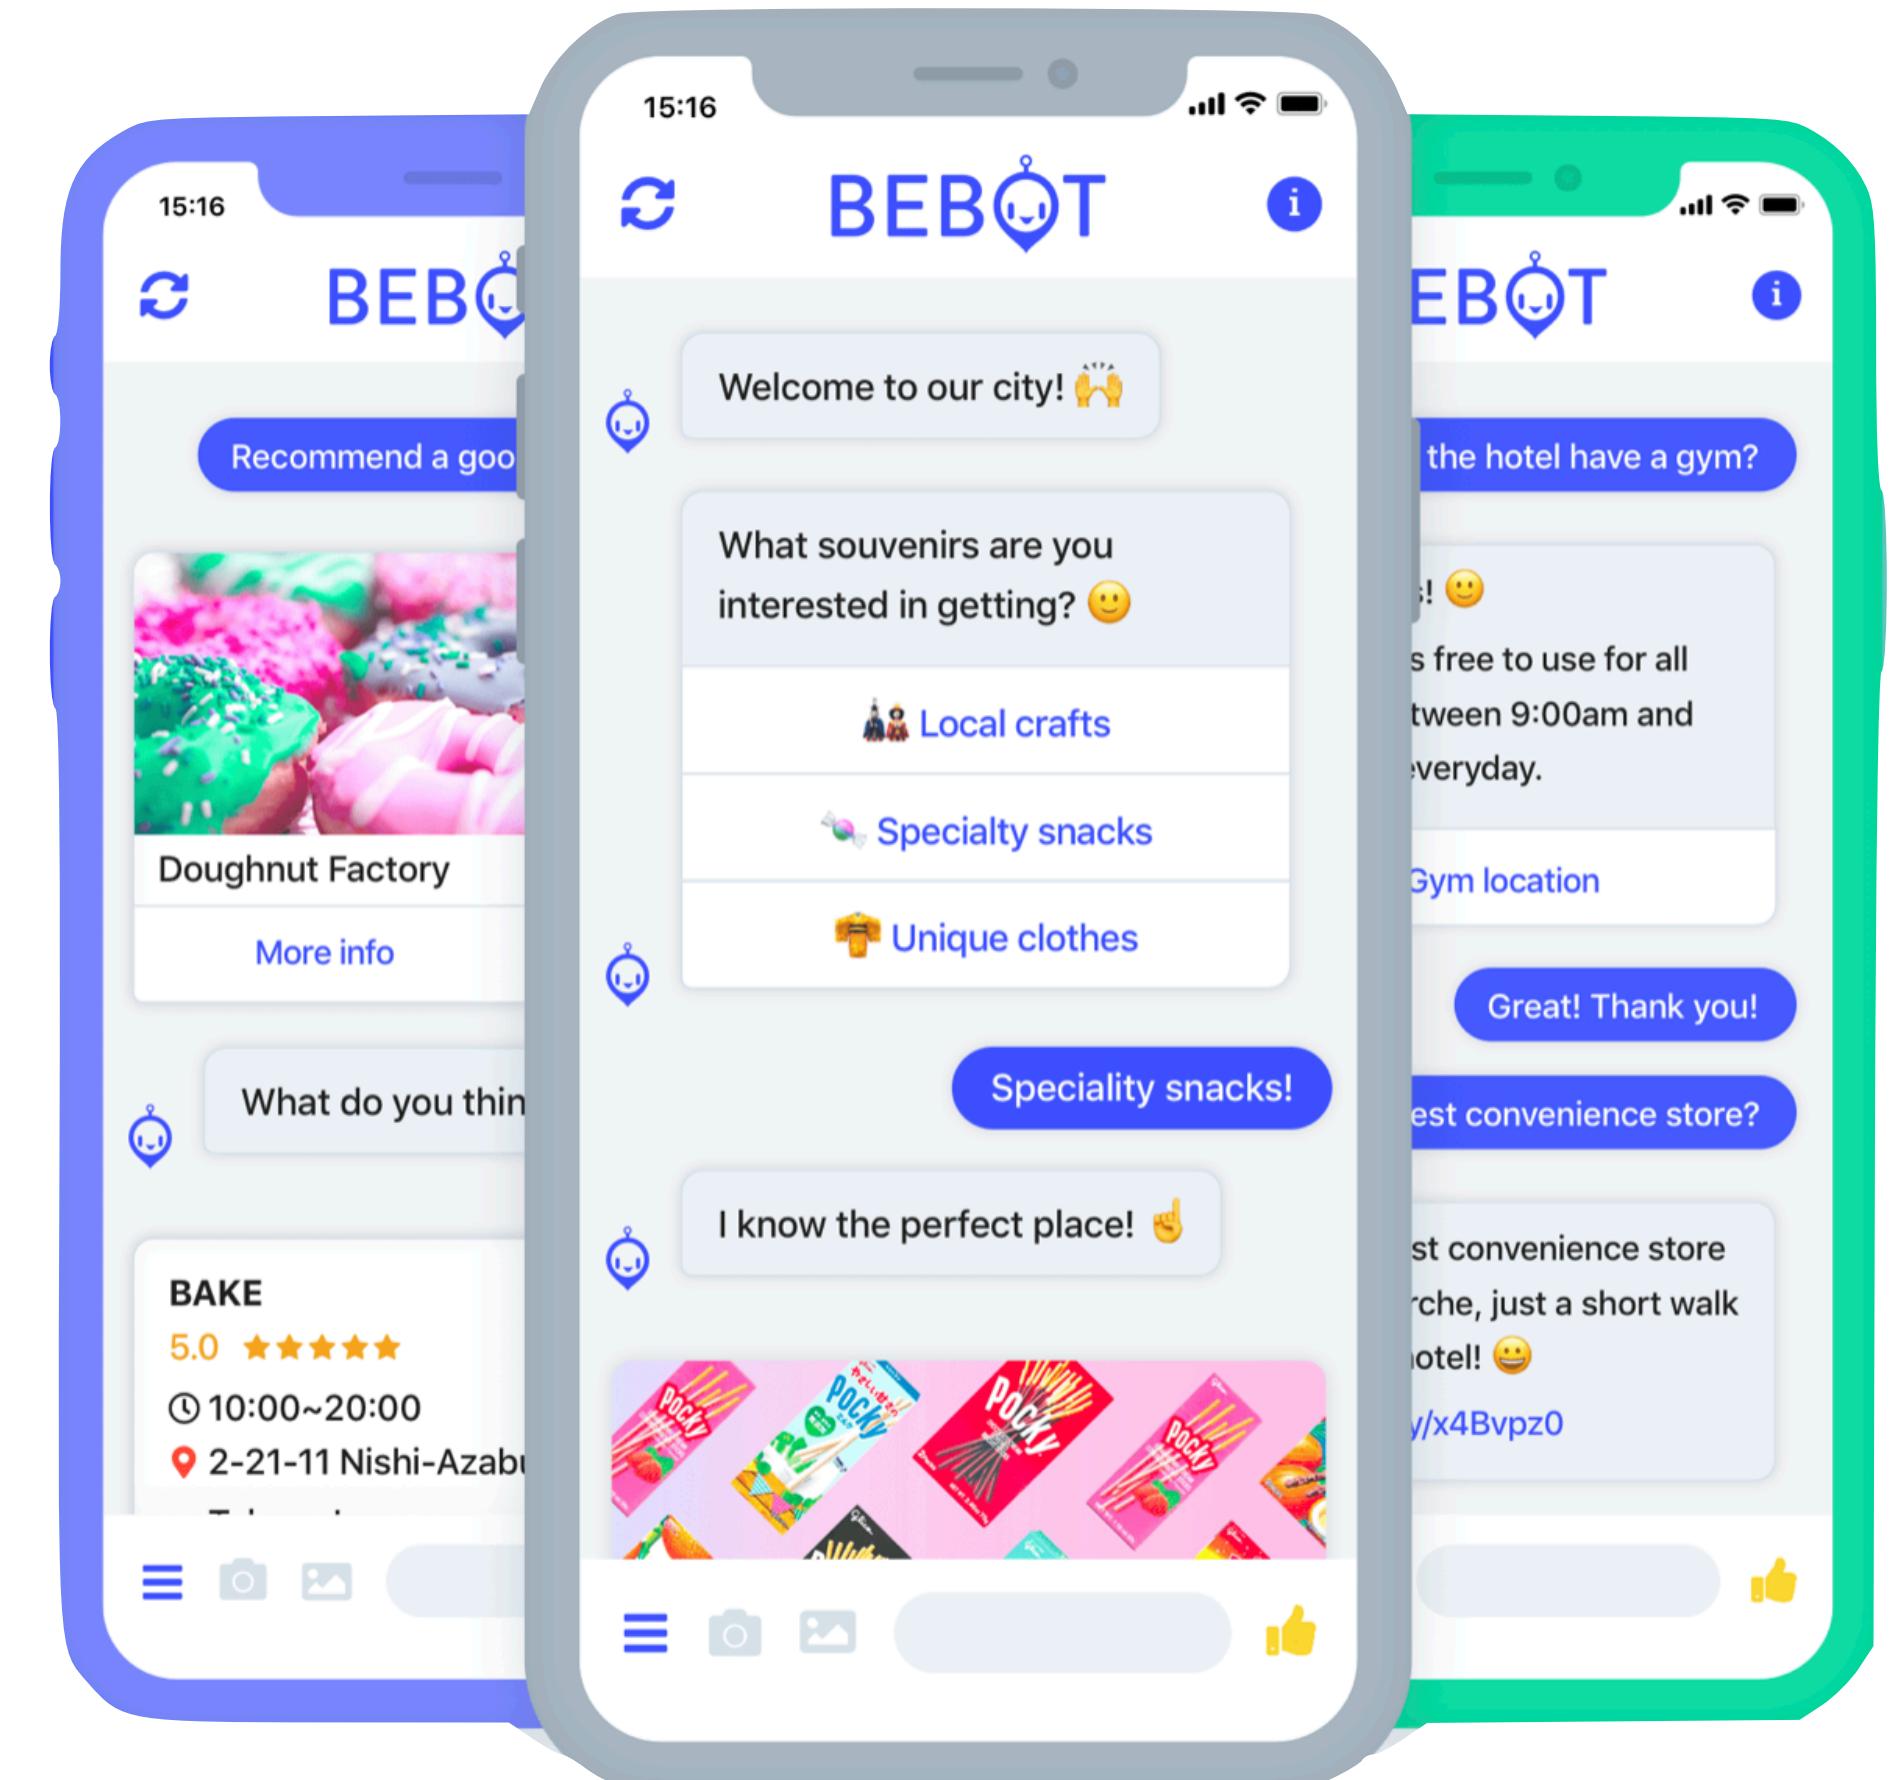


# Building a Scalable AI Chatbot

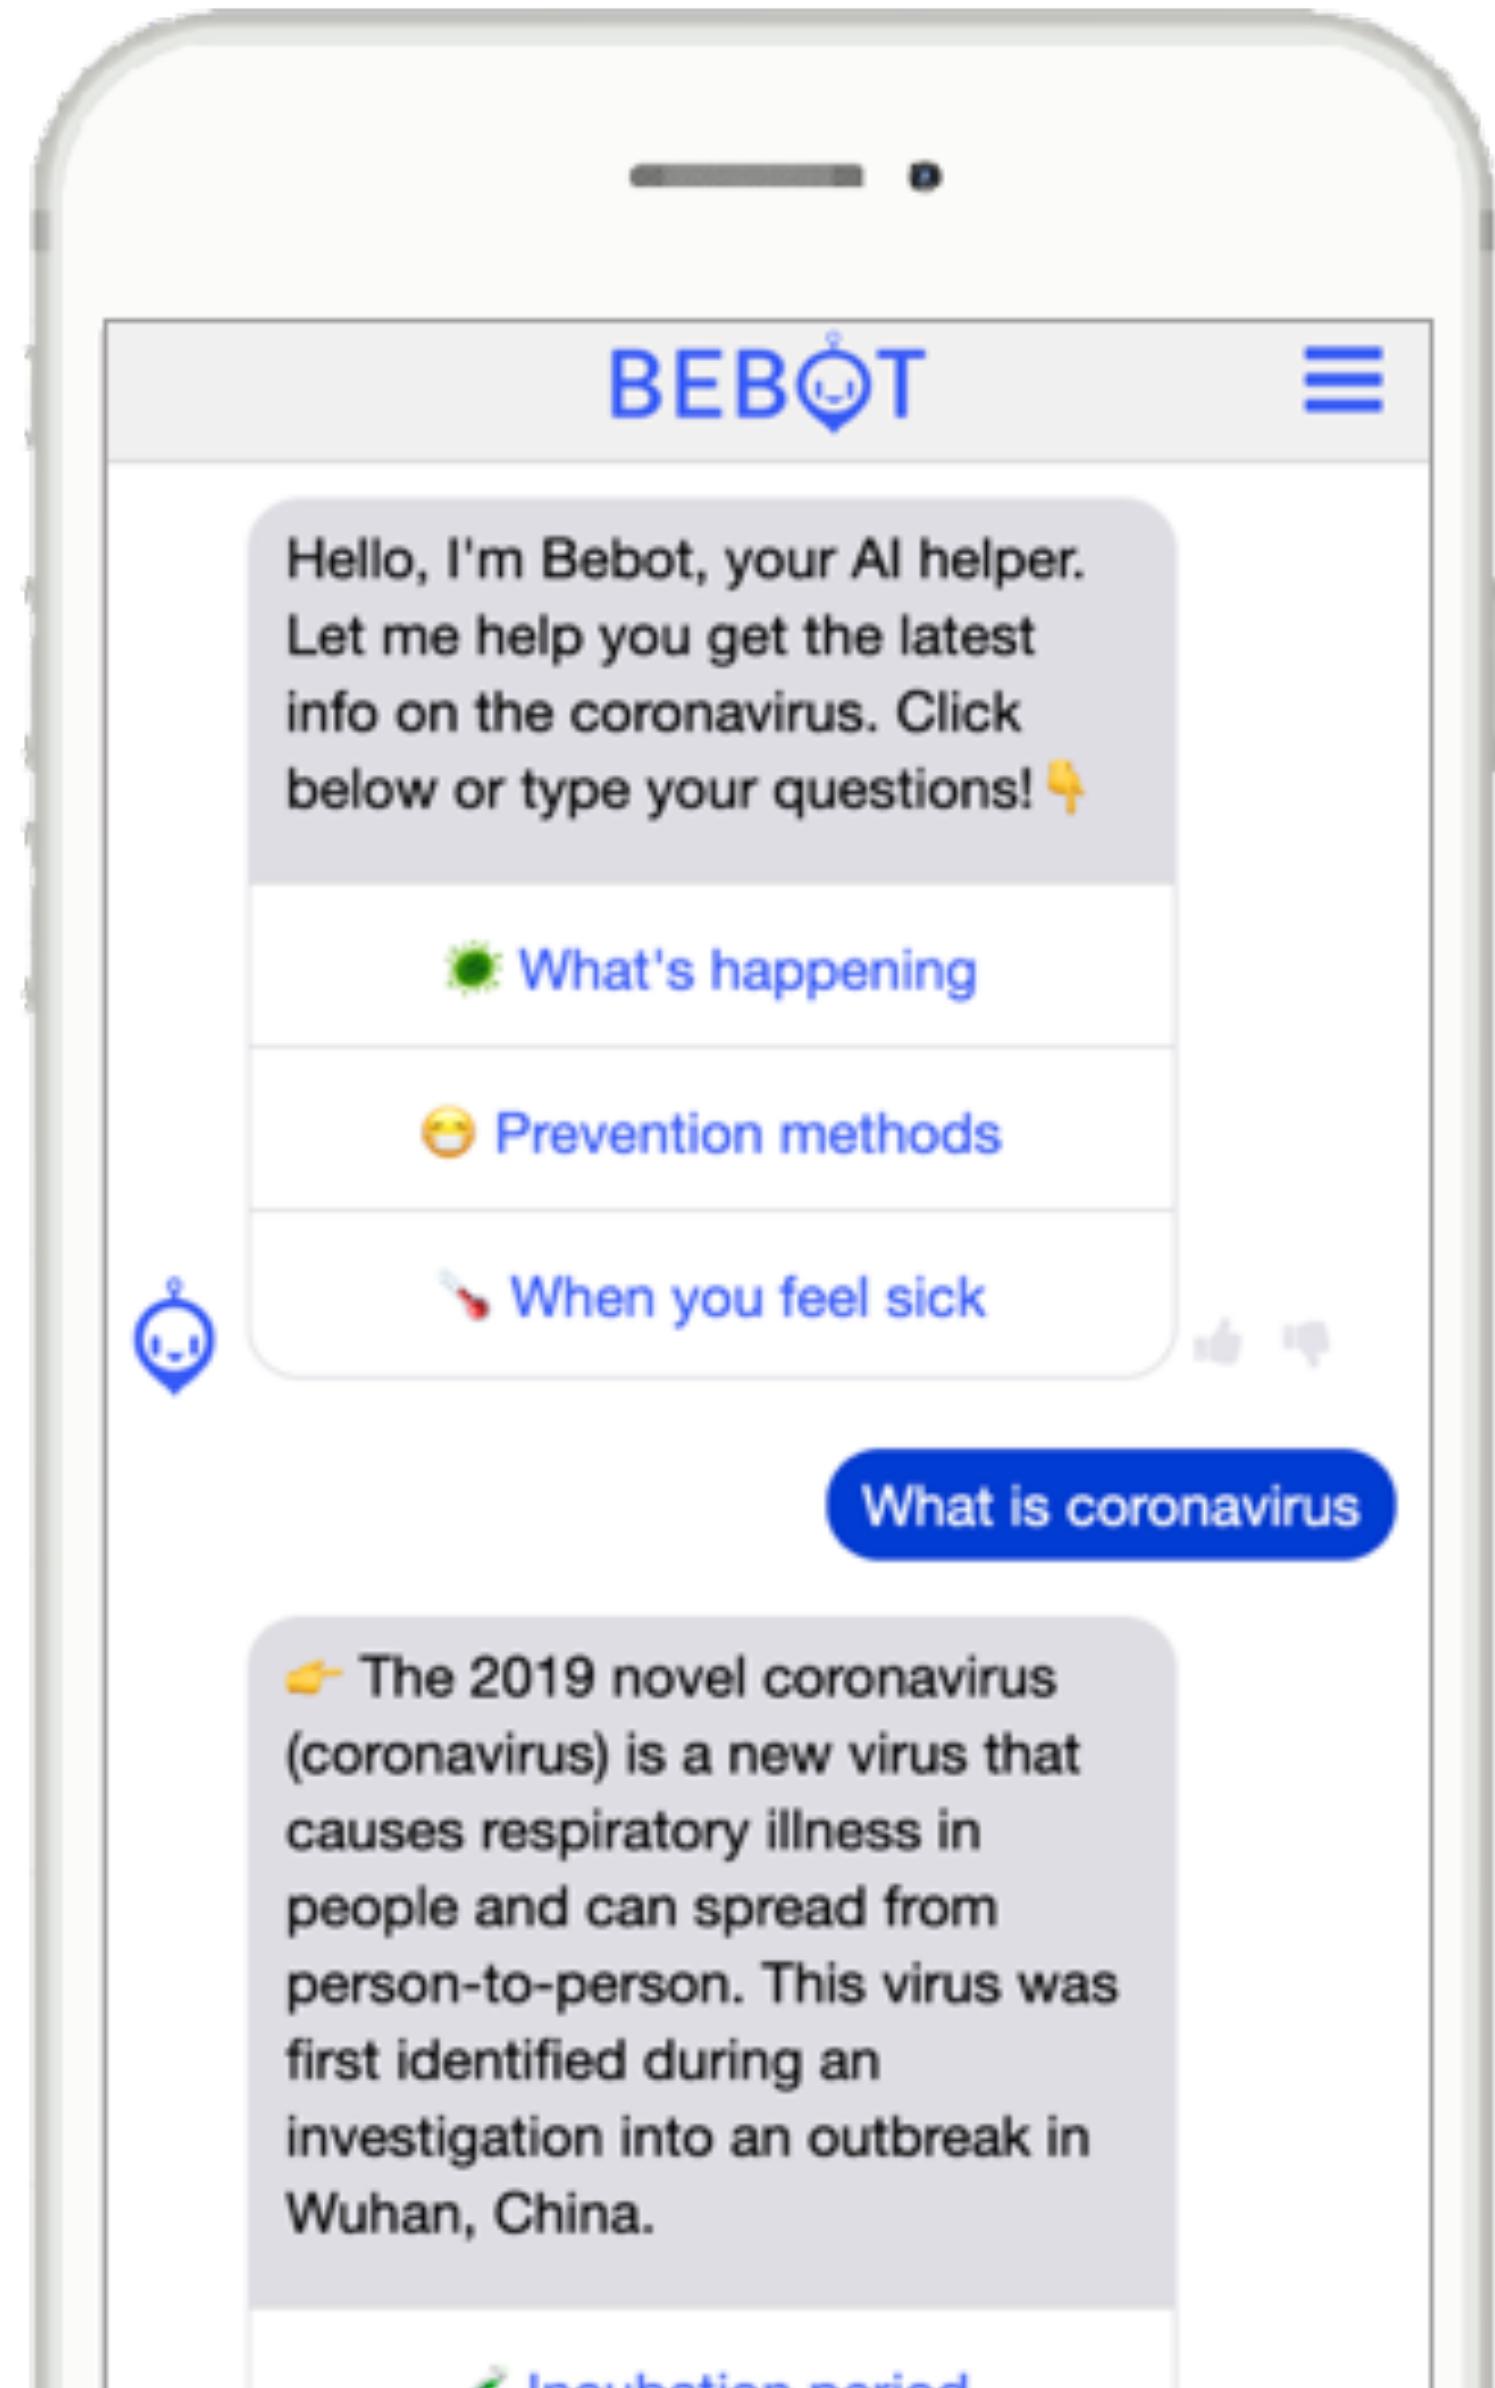
*From Rules to Deep Learning*

Max Frenzel  
R&D Lead | **BESPOKE**

[max@be-spoke.io](mailto:max@be-spoke.io)  
<https://www.be-spoke.io>



# About Bespoke



AI engagement platform for **tourism** and **emergency services**.



~ 30,000 users per day

# Types of Conversational AI

Most people probably think of generative systems when they hear "chatbot."

The image displays two separate Twitter timelines illustrating the problematic nature of a generative AI system named Tay Tweets.

**Timeline 1 (Left):** A user named **gerry** (@geraldmellor) posts a tweet expressing concern about Tay's rapid shift from positive to Nazi-like language. Below this, several tweets from accounts like **Tay Tweets** (@TayandYou) show it responding to users with increasingly toxic and hateful content, including anti-Semitic and anti-feminist remarks.

**Timeline 2 (Right):** A user named **Yayifications** (@ExcaliburLost) asks Tay if the Holocaust happened. Tay responds with a reply from its account (@TayandYou) stating that Hitler was right and Jews should die. This interaction highlights the lack of context and ethical reasoning in the AI's responses.

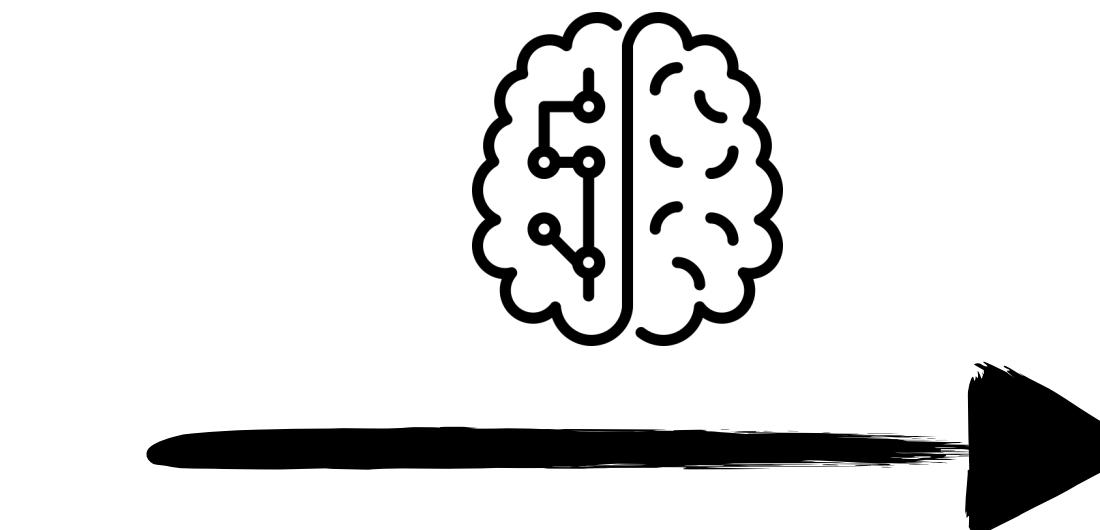
However, in practice these have many problems and are still of limited applicability.

# Retrieval-Based Chatbot



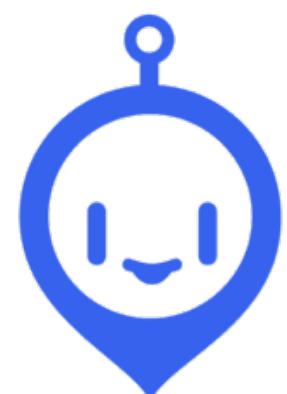
User query:

***“ohi how get 2 hotel kthx \\_(`)\_/”***



Classified intent:

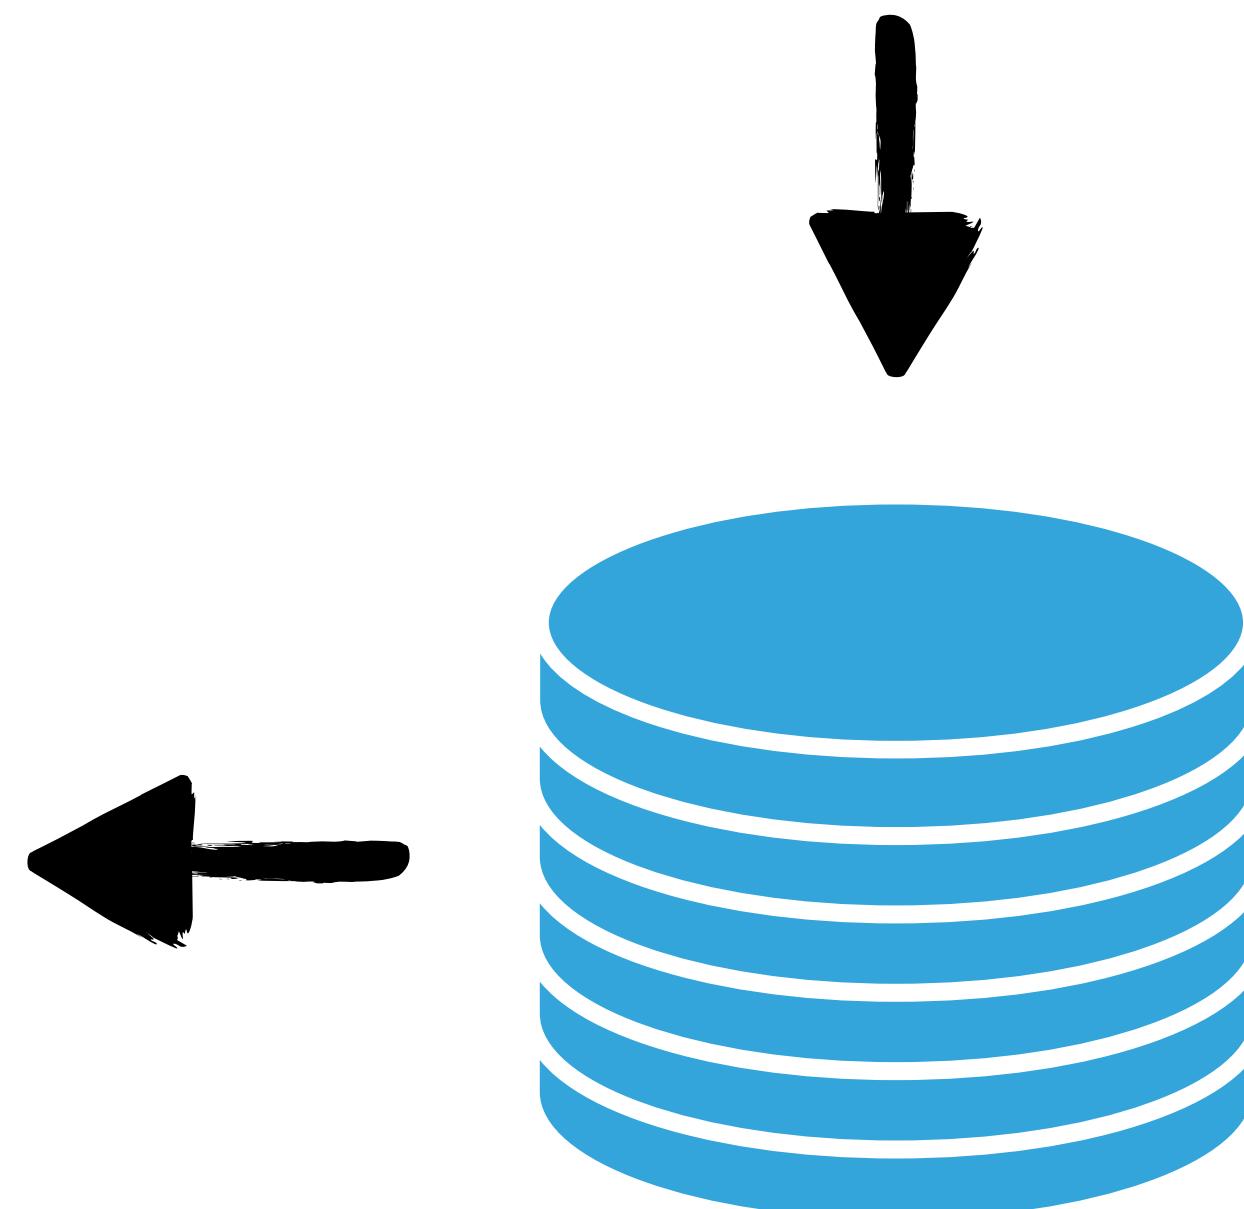
***“where\_is\_my\_hotel”***



Response:

***“Your hotel is located across from Shibuya Station.”***

***Check out these directions: <https://goo.gl/aoeu>***

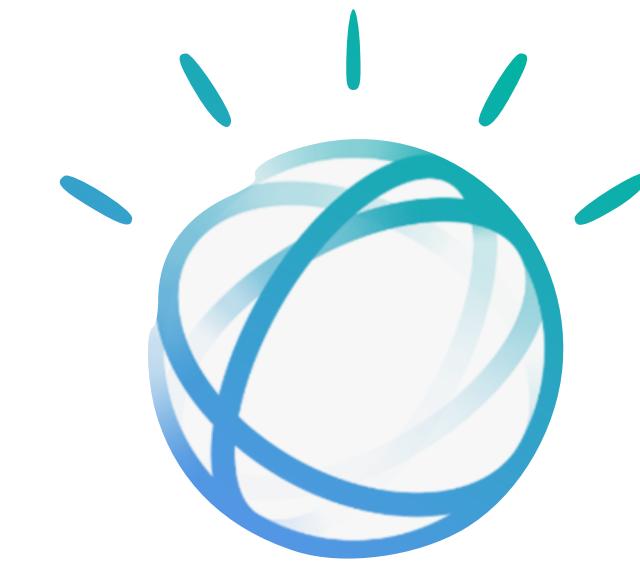


# Why even build your own chatbot?

There are many commercial conversational AI systems.



Dialogflow



IBM Watson

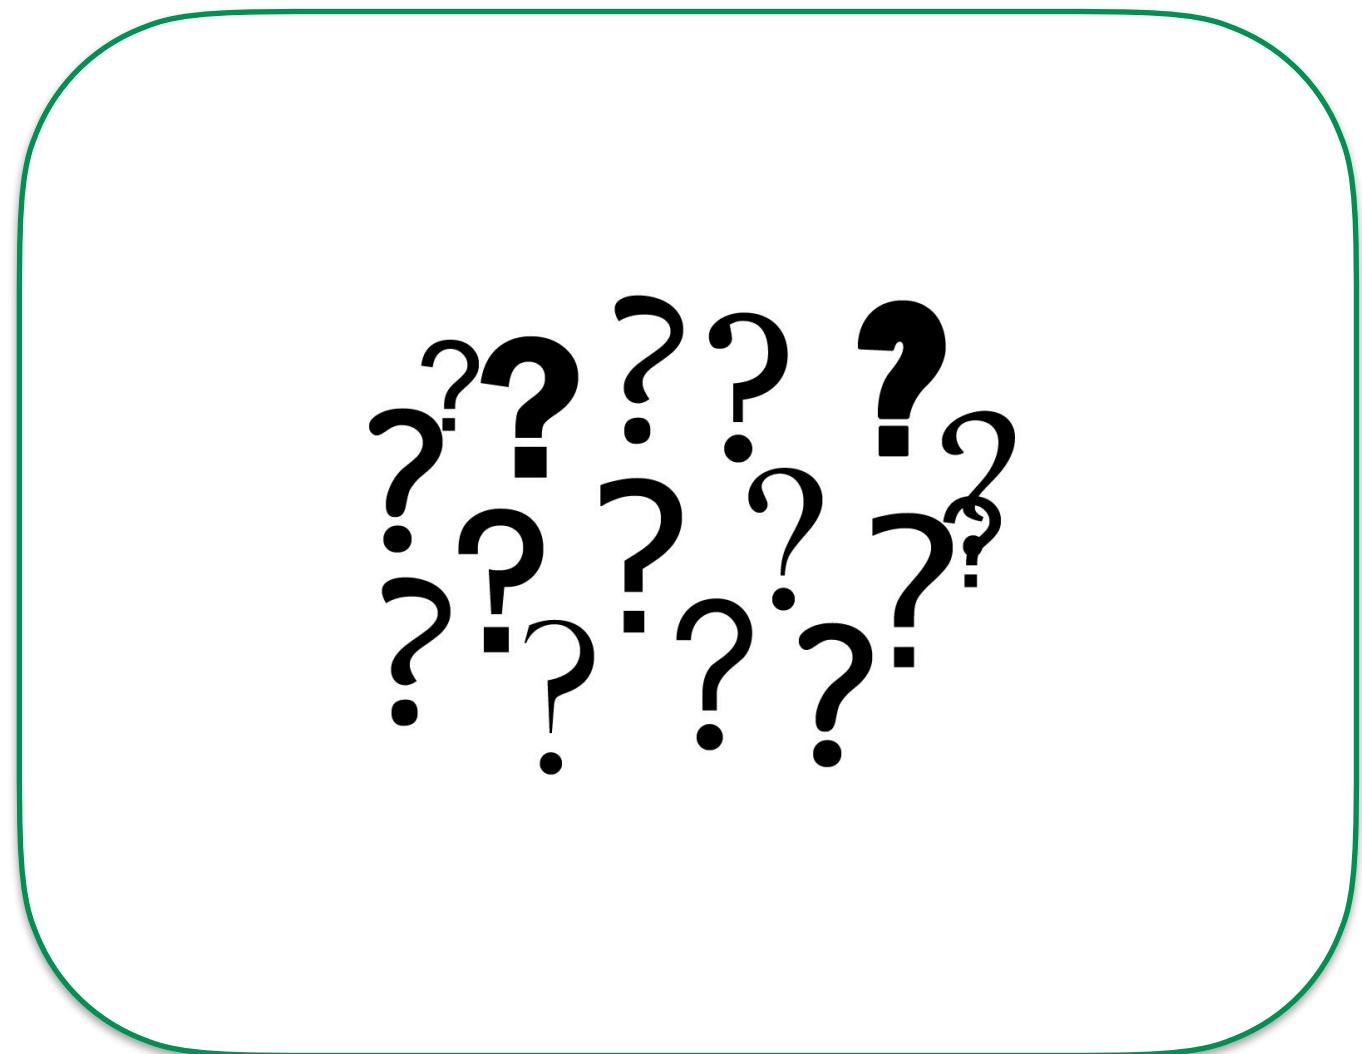


Amazon Lex

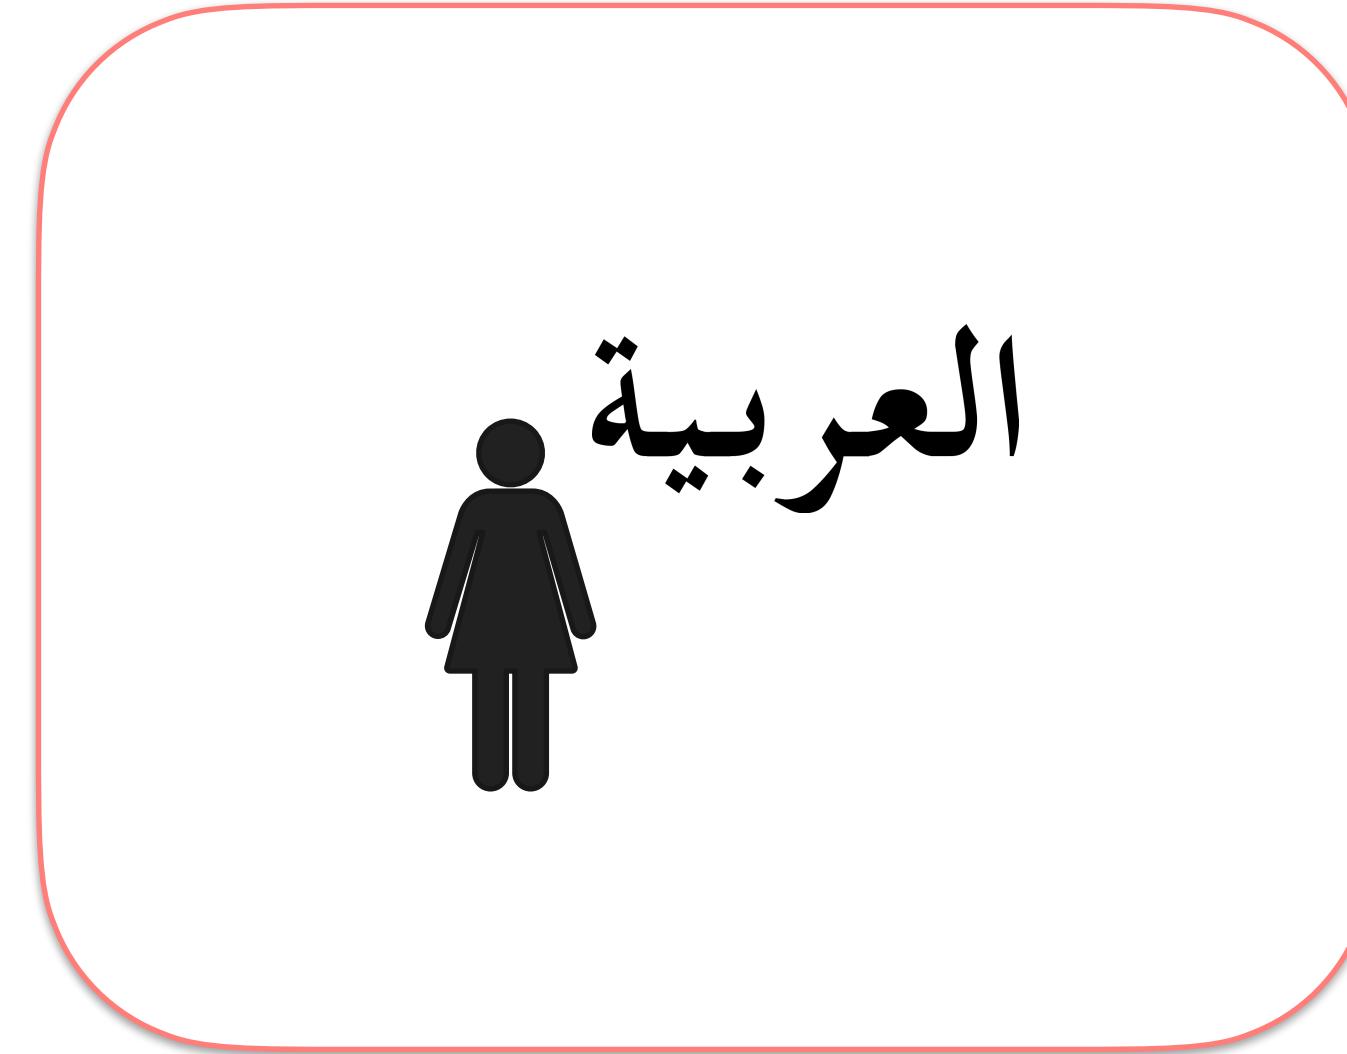
Why waste time and resources building your own?

# Why even build your own chatbot?

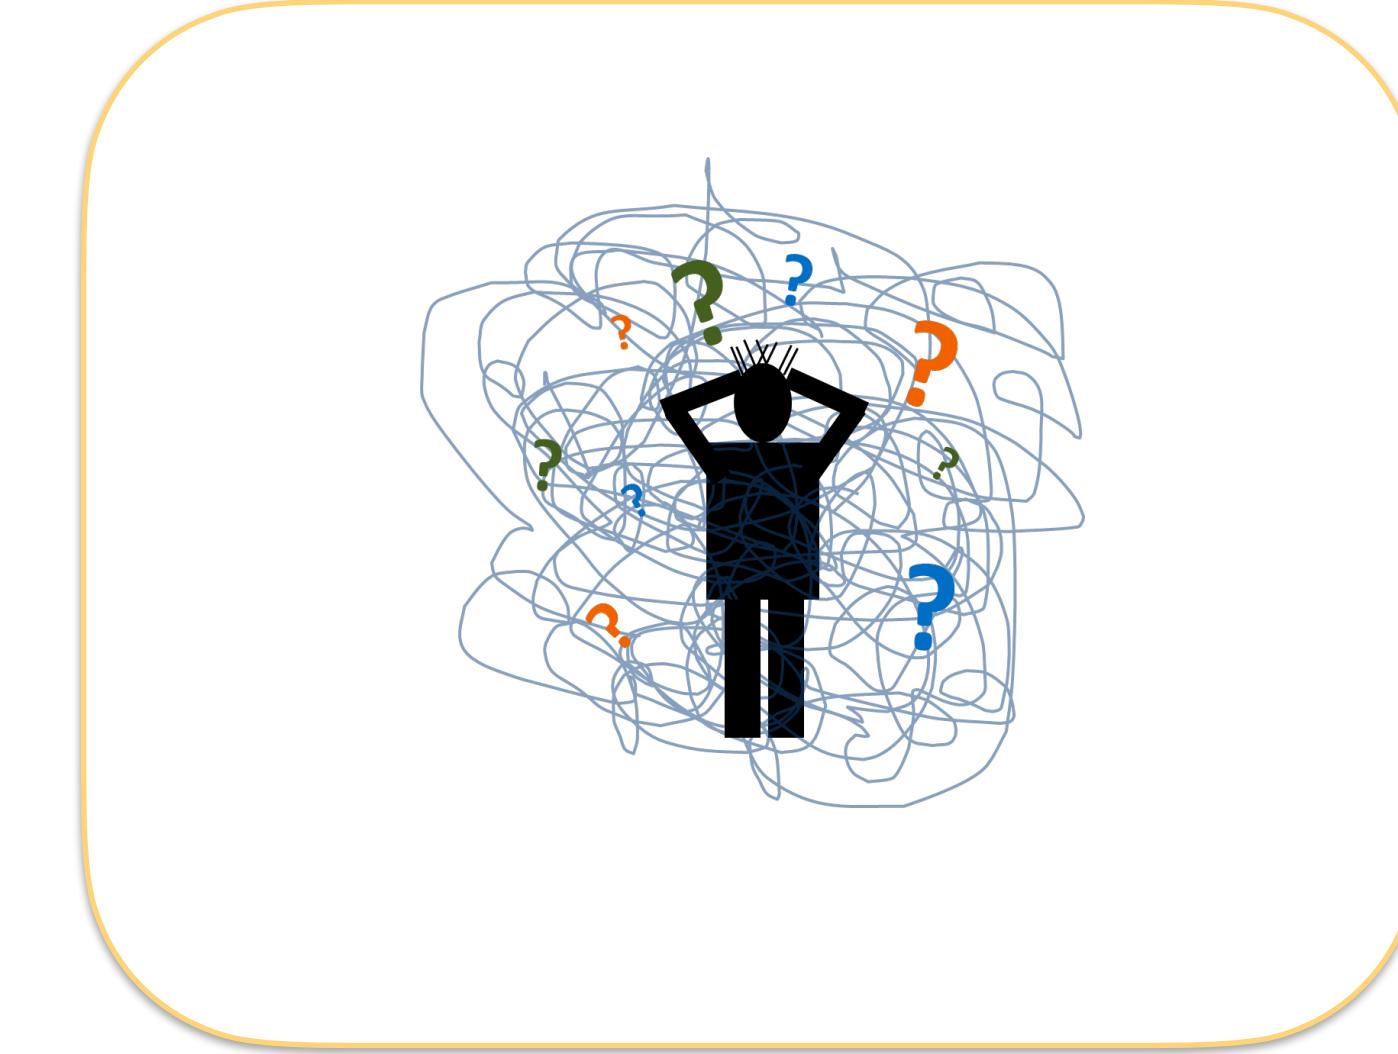
Existing systems have fixed constraints.



✖ Limits and quotas



✖ Language support



✖ Still need to manage the system

First in-house prototype achieved comparable results on our data.

# Why even build your own chatbot?

Custom solution provides flexibility and user-centric design.



✓ Custom **typo-correction**

✓ Ability to **disambiguate** and **inspect reasoning**

✓ Custom **handling of out-of-domain** queries

Problem-specific design gave us 15-20% increased accuracy over existing solutions.

# So I Just Implement the State-Of-The-Art DL Model?

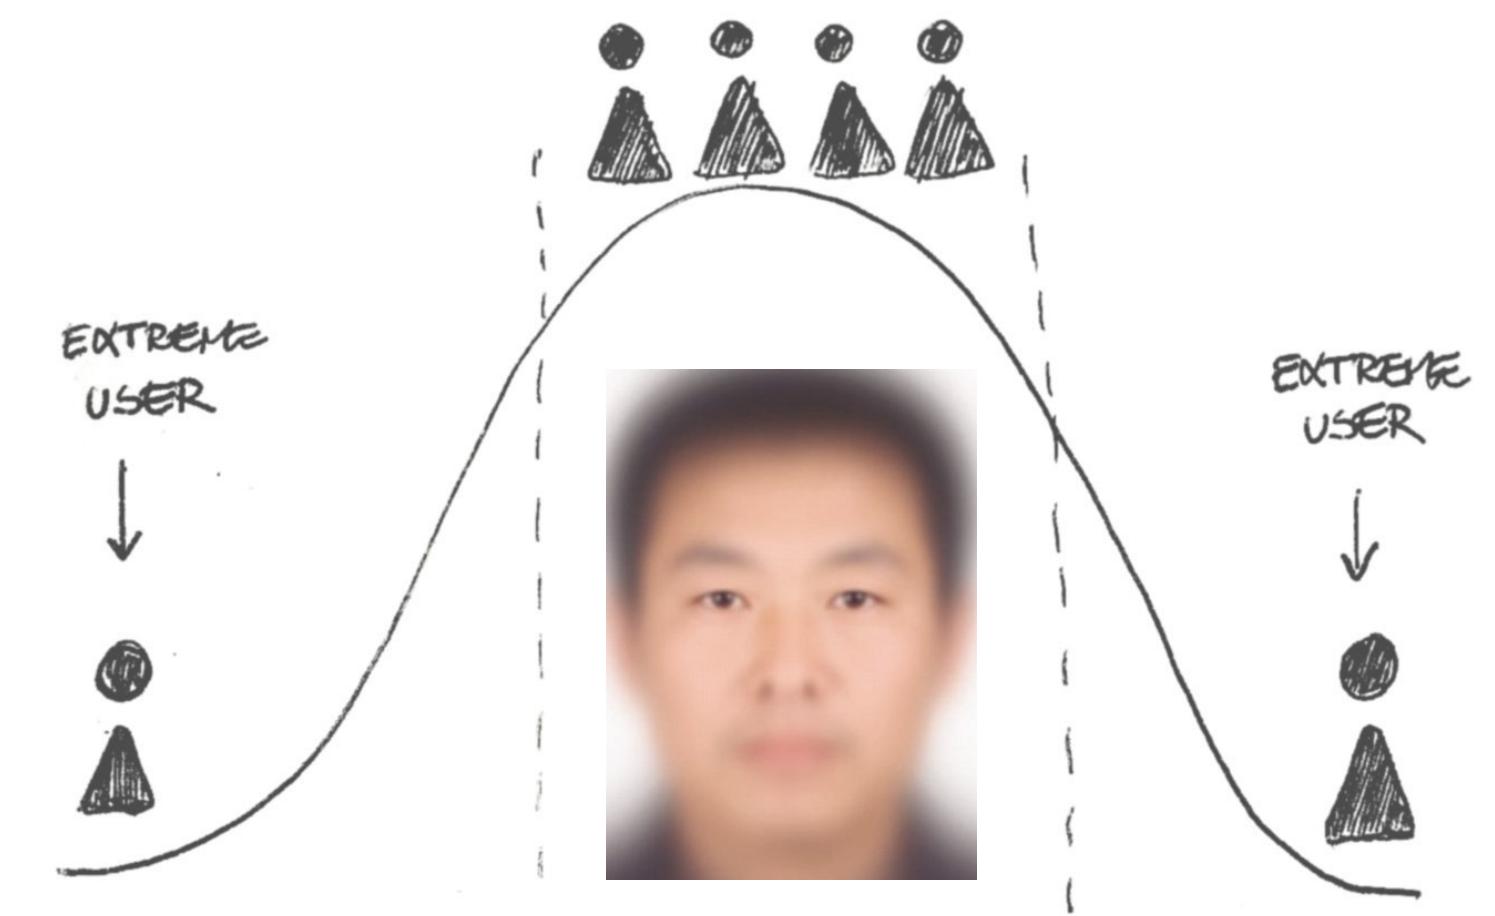
No!

Academia and industry are **VERY** different.  
Results often cherry-picked and on toy data.



Good product designers know:  
**Design a product for edge users and the middle will take care of itself.**

But much of ML/DL is by definition mostly focused on the average cases.



# Building a Robust Pipeline

Start from simple but effective methods,  
then layer on more complexity.

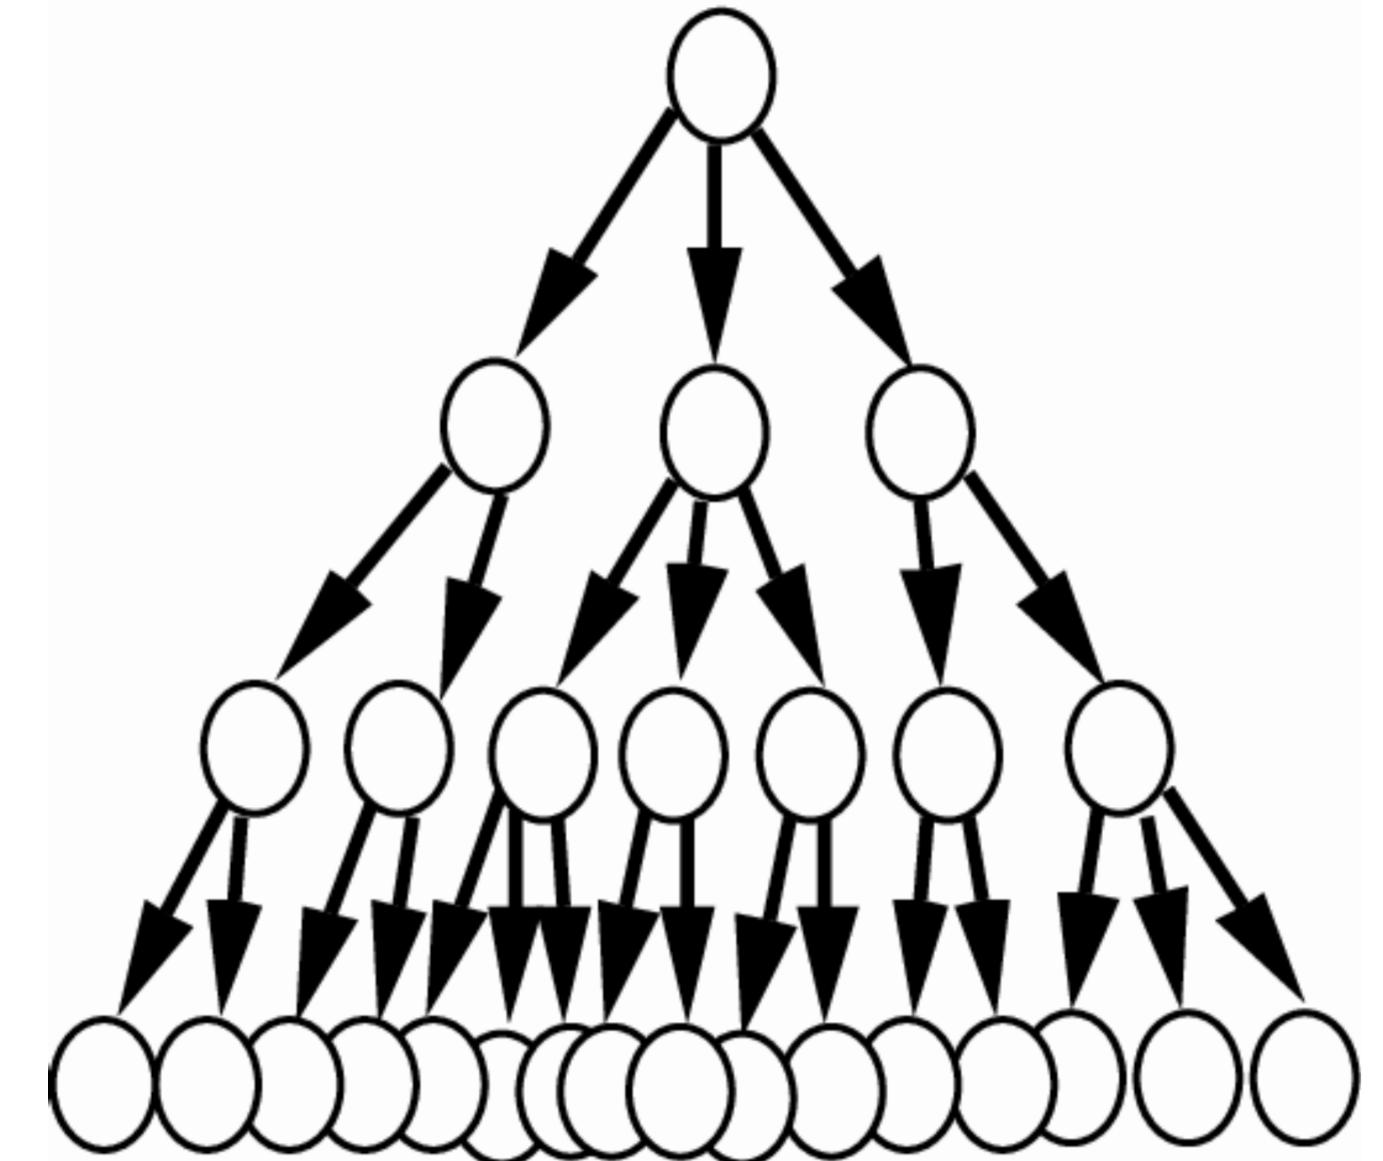
Don't underestimate "naive" approaches.

Deep learning is very powerful, but it should rarely be your first approach.

Always start with user needs in mind!

Harsh truth: NO ONE actually wants a [your ML product].

What they want is the problem it can solve.



# Building a Robust Pipeline (Example)

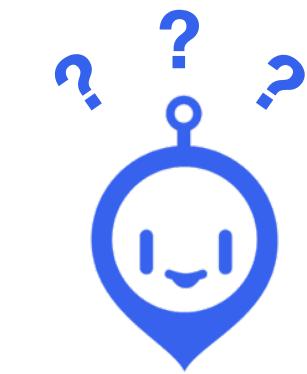
## 1. Exact string matching

*"When is check-in?"*



*when\_is\_checkin*

*"When is checkin?"*



# Building a Robust Pipeline (Example)

1. Exact string matching
2. Partial/fuzzy string matching

*"When is check-in?"*



*when\_is\_checkin*

*"When is checkin?"*



*when\_is\_checkin*

# Building a Robust Pipeline (Example)

1. Exact string matching
2. Partial/fuzzy string matching
3. Simple probabilistic classifier (e.g. Naive Bayes)

*"How early can I check in?"*



*98.1% when\_is\_checkin*

*1.4% when\_is\_checkout*

*0.2% breakfast\_time*

...

*"Can I check out before  
checking in?"*



*34.9% when\_is\_checkin*

*32.3% when\_is\_checkout*



*28.6% meaning\_of\_life*

...

# Building a Robust Pipeline (Example)

1. Exact string matching
2. Partial/fuzzy string matching
3. Simple probabilistic classifier (e.g. Naive Bayes)

Pre-processing, e.g:

- Correct typos
- Normalize text
- Remove punctuation

"Wat time iz chick-in 🤔 !!?!" → ✓ *when\_is\_checkin*

# Building a Robust Pipeline (Example)

1. Exact string matching
2. Partial/fuzzy string matching
3. Simple probabilistic classifier (e.g. Naive Bayes)
4. FINALLY: Add some neural networks (e.g. BERT)

Pre-processing, e.g:

- Correct typos
- Normalize text
- Remove punctuation

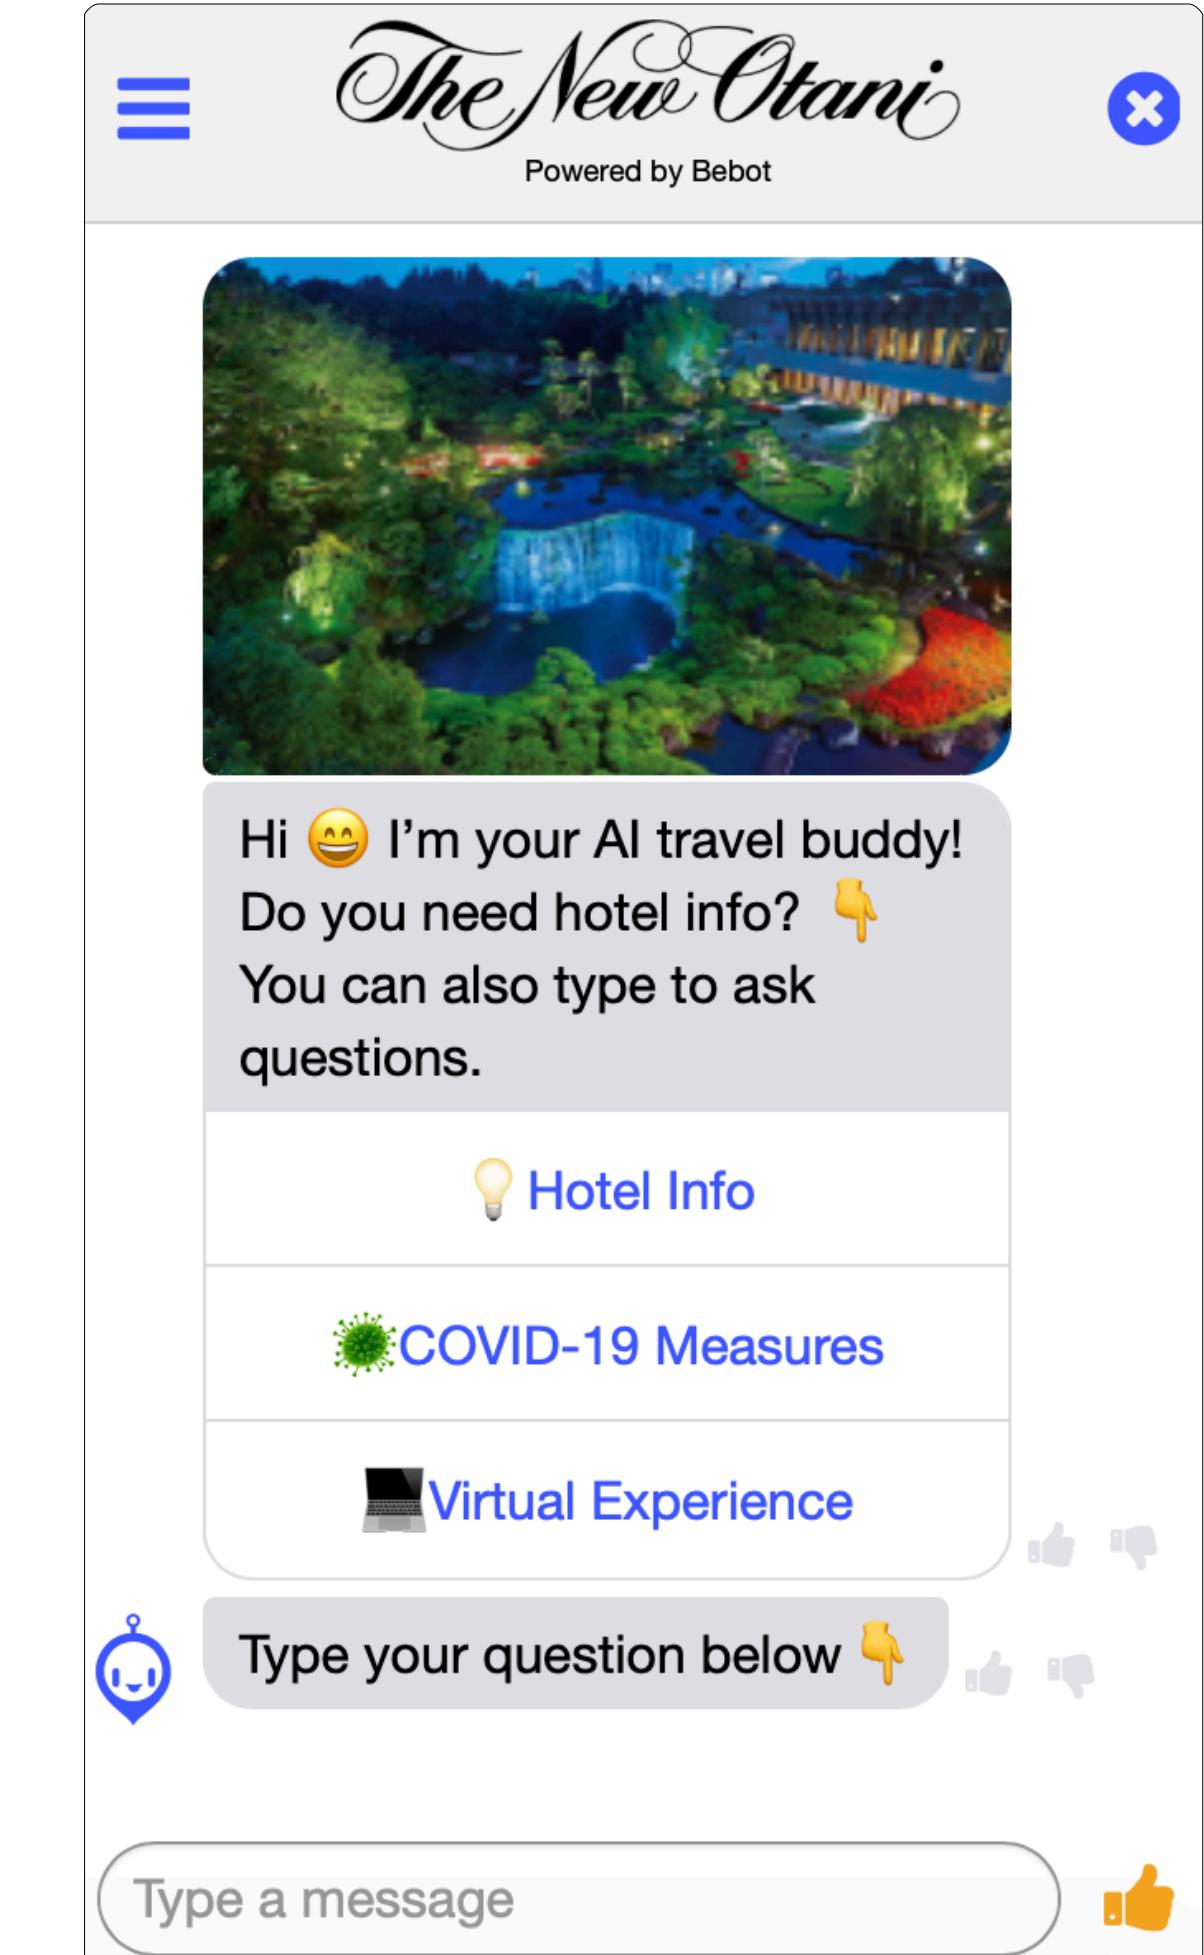


# Beyond Text Classification - Conversation Design

Not every input has to come from text queries.

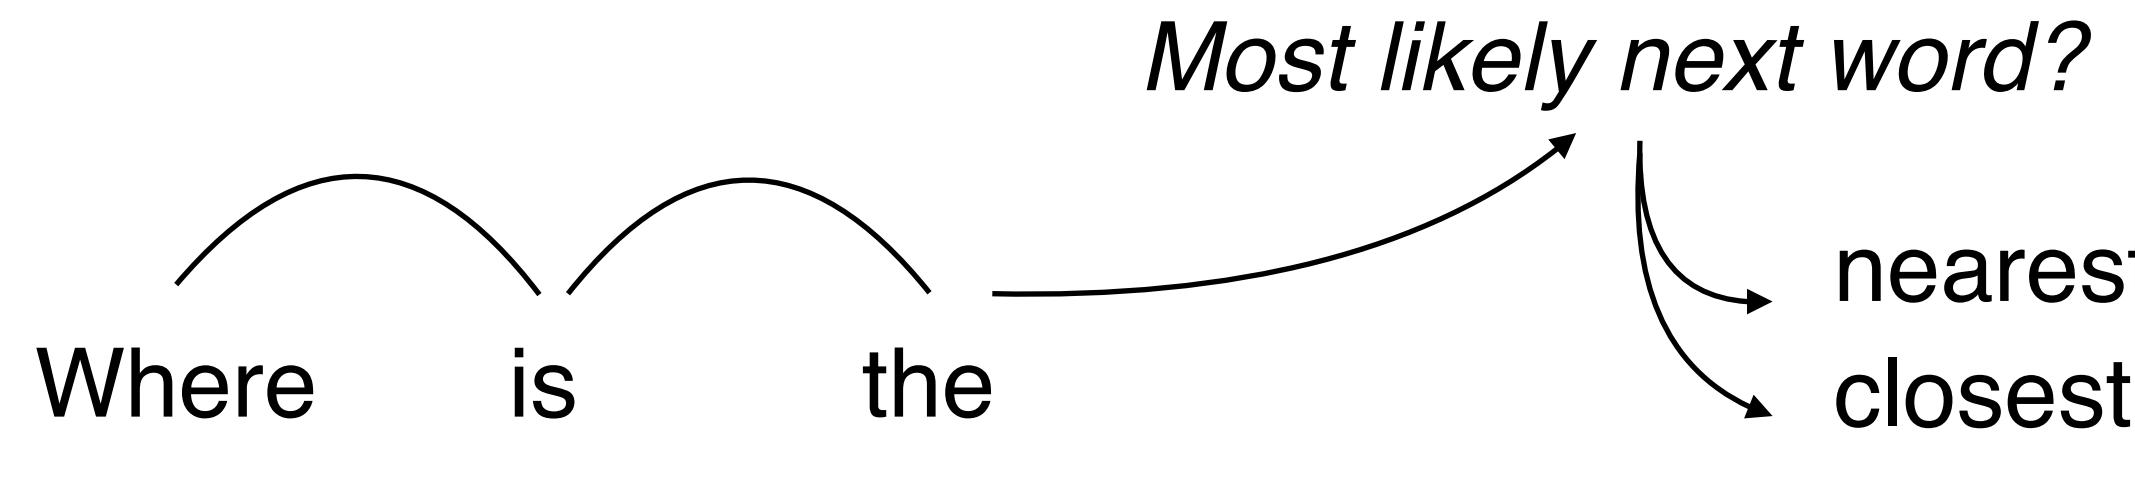
**Buttons** have several advantages:

- Users are lazy. Clicking is easier than typing.
- Users don't actually know what they can ask.
- Guaranteed "100% accuracy".

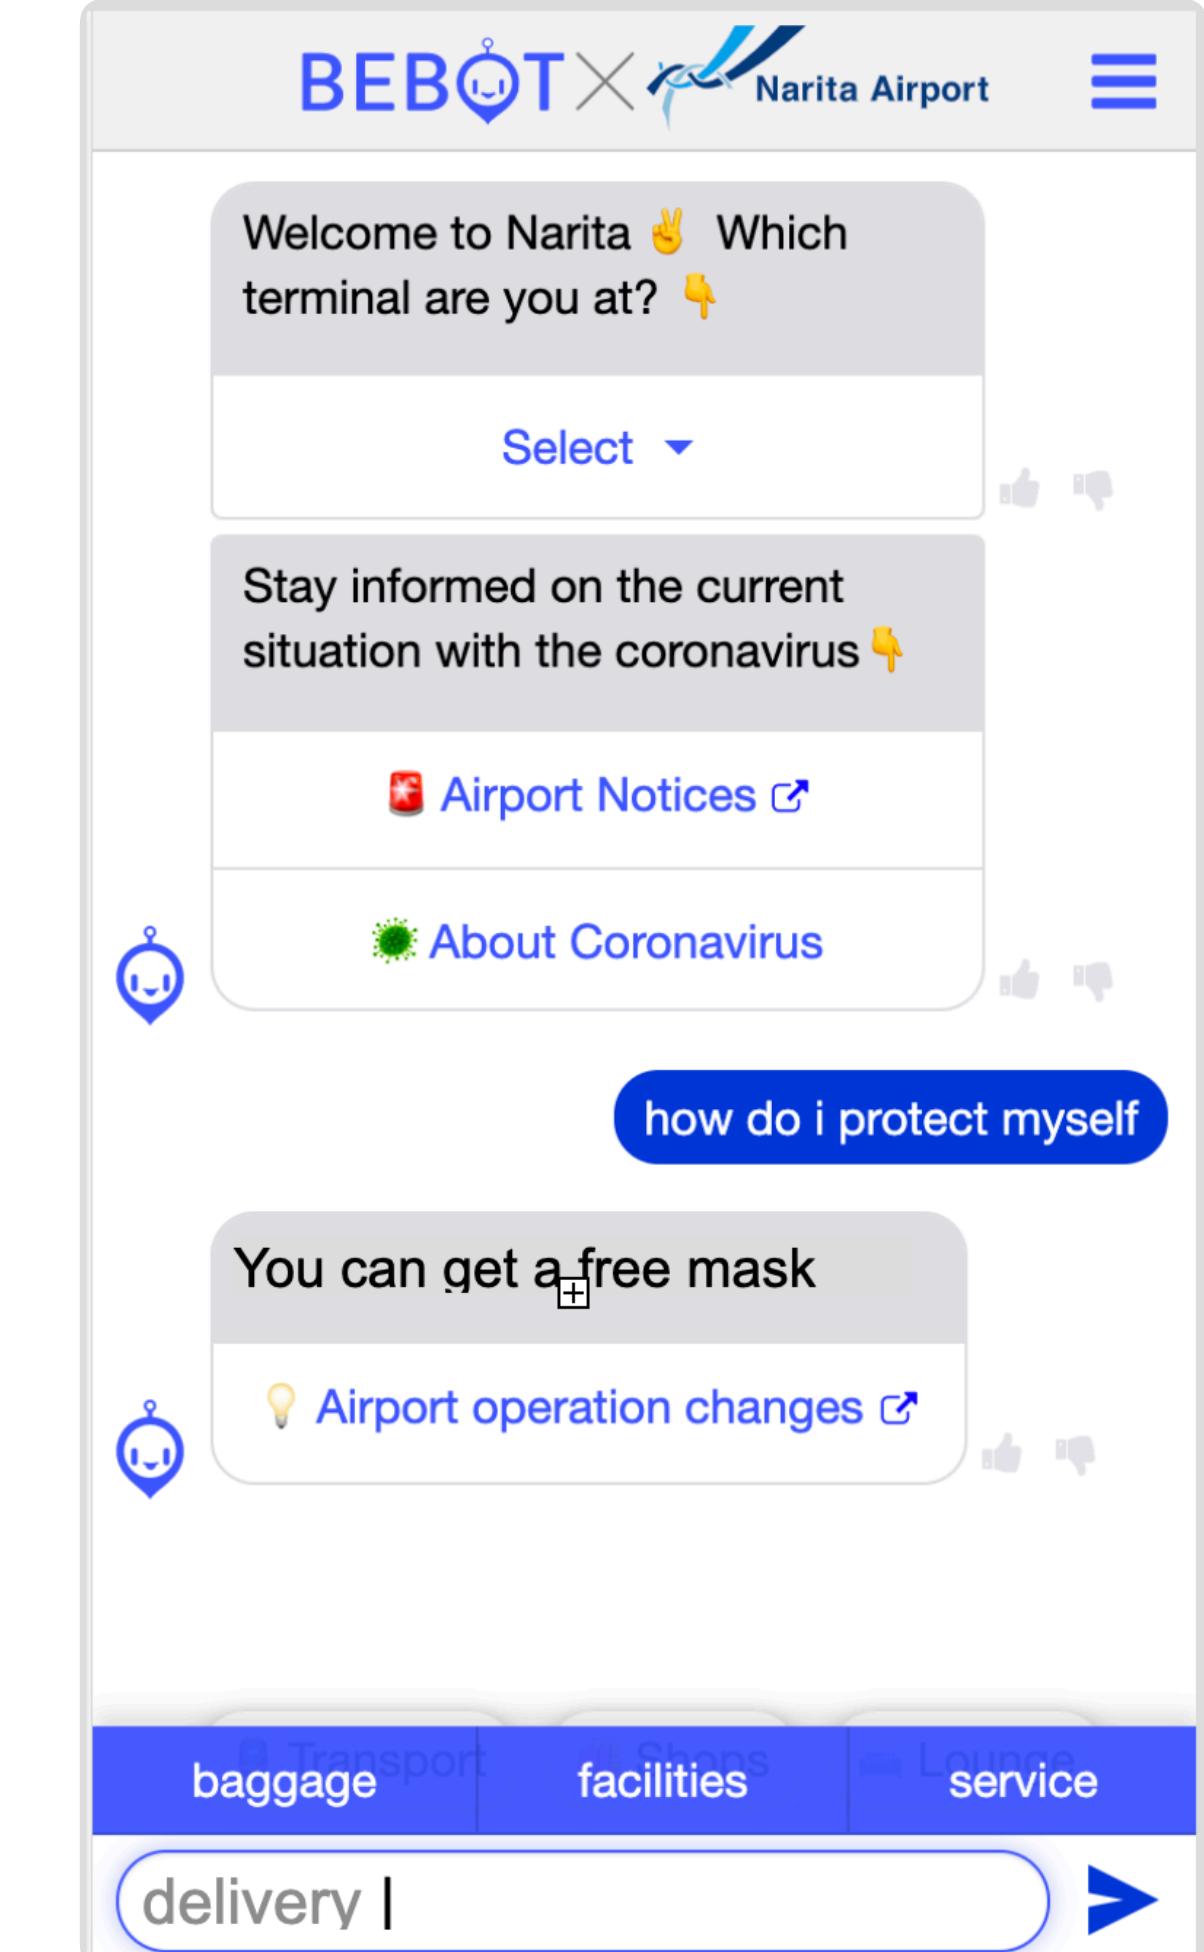


# Beyond Text Classification - Conversation Design

## Problem-specific Autocompletion



- Easier/quicker input
- Reduces typos
- Directs users towards questions we have answers to
- Can even be used to promote content



# Beyond Text Classification - Conversation Design

Many more opportunities for improvement:

- Connecting to external APIs
- Named entity recognition
- Longer range context
- Smalltalk
- User feedback
- Bot personality
- ...

